

CoQuí scGW on GPU: Performance Diagnosis and Optimization Plan

Why the port runs at CPU speed — and the path to 5–8×

Miguel Morales (with Claude)

Flatiron Institute — rusty cluster — branch `gpu`, `ad9bc48`→`765754c`

July 26, 2026

Abstract

We benchmarked CoQuí’s THC-based self-consistent GW on rusty CPU (Genoa) and GPU (A100, H100) nodes using Si supercell problems up to $\Pi \approx 1.4$ TB ($4 \times 4 \times 4$ k -grid, 500 bands). The GPU compute kernels are excellent — the screened-interaction Dyson solve runs 16–23× faster than the CPU code and sits at its roofline — but end-to-end the port shows essentially no speedup (1.0–1.5× at equal node count). The cause is quantified precisely: $\sim 70\%$ of GPU wall time is data movement — host-staged all-to-all redistributions in the $\tau \leftrightarrow \omega$ transforms, a host-resident gemm in the THC basis transform, per-iteration re-gathers of iteration-invariant data, raw `cudaMalloc` churn, and host-only SCF phases. A single-node (network-free) GPU run proves the cost is PCIe staging, not the interconnect. No algorithmic change is required: the existing transpose-based transform is communication-optimal. We propose a prioritized engineering plan (device-direct chunked NCCL redistribution, device memory pool, device-resident basis transforms, invariant caching, device Dyson stage) with a projected $\approx 5\times$ node-for-node speedup at the $2 \times 2 \times 2/500$ -band scale, growing to 5–8× at production scale. During the study we also found and fixed a correctness regression that had made the THC-HF Fock matrix processor-count-dependent on *both* CPU and GPU since 2026-05-12; all configurations now agree to 10^{-13} .

1 Methodology

All runs use one source tree (branch `gpu`, HEAD `765754c`) and three binaries: CPU (icelake SLATE), CUDA sm80 (A100), CUDA sm90 (H100/H200). Identical physics inputs throughout: $\beta = 2000$, $\Lambda = 1200$, IR grid “high” ($n_\tau = n_\omega = 137$), THC `ecut`=150, `thresh`= 10^{-5} , damping mixing 0.7; three SCF iterations (two at kp444). Phase timers are device-synchronized. Because THC interpolating-point selection depends on the processor grid (finding R1), all cross-configuration comparisons load a single pre-built ERI file (`save`= mechanism). Every timed configuration was validated against a serial CPU reference energy.

Node type	Configuration	FP64 DGEMM/node	Host RAM
CPU (ccq, Genoa)	$2 \times 48c$ AMD Genoa, 96 ranks	≈ 4.5 TF	1.5 TB
GPU A100	$4 \times$ A100-SXM4-80GB (NVLink)	39 TF vec / ≈ 78 TF TC	1 TB
GPU H100	$8 \times$ H100-PCIe-80GB	≈ 270 TF TC	1 TB

Table 1: Hardware. Fair yardstick: gemm-dominated phases should run ≈ 10 – $17\times$ faster on an A100 node than a Genoa node (≈ 25 – $60\times$ on $8 \times$ H100), minus communication.

2 Correctness findings

Timing data is only meaningful on correct code; the study began by validating physics and found the following.

#	Status	Defect	Impact
B1	fixed (765754c)	<code>_aux_to_primary_impl</code> X-factor cache (introduced 734ae79, May 12) reused the P -sliced cache as the right gemm factor when $ip=iq$; the right factor must be sliced with <code>Q_rng</code> .	THC-HF Fock processor-count dependent on CPU <i>and</i> GPU since May 12 ($\ \Delta F\ _\infty$ up to 0.3 a.u. between rank counts; GPU total energies off $\times 1000$). GW Σ provably unaffected (its grid keeps $np_P=np_Q=1$; bit-identical across ranks). All May benchmark energies were wrong. Post-fix: CPU 1/2/4/8 ranks and GPU A100(16r)/H100(1–16r) agree to 10^{-13} .
B2	fixed (80f4a5d)	<code>TimerManager</code> aborts in <code>update_w<HOST></code> : the timer dump queries DEVICE-only timer names.	Every freshly built CPU binary crashed in iteration 1.
B3	open / workaround in place	Device ISDF/THC builder seg-faults at ≥ 16 ranks: UCX “invalid permissions for mapped object” in the <code>Zqur</code> stage (device buffer reaches a host-only MPI path; forcing UCX CUDA transports does not help). Pre-existing (May 16-rank run died identically).	Blocks all large GPU runs. Workaround: pre-build the ERI at ≤ 8 GPU ranks or on CPU, <code>save=</code> , load everywhere (also guarantees identical integrals across runs).
B4	open	μ -search bracketing walk (<code>scf_common.cpp:190–202</code>) is silent and uncapped (0.2 a.u. steps, no iteration limit, no logging).	Turned B1’s wrong Σ into a 4 h wall-limit hang. Needs a cap + diagnostics.
B5	open	<code>compute_eigenspectra</code> allocates the full $\Sigma(\omega)$ ($n_\omega n_s n_k n_b^2$) <i>per rank</i> : 35 GB/rank at kp444/500b.	kp444 CPU jobs OOM at full rank density (forced to 24 ranks/node). Distribute or chunk over ω .
R1	reproducibility	THC interpolating-point selection depends on the processor grid ($N_p = 4482/4483/4471$ across rank counts).	$\sim 10^{-4}$ a.u. total-energy spread between rank counts at thresh 10^{-5} . Comparisons must share a saved ERI.

Table 2: Defects found during the study.

3 Measured performance (corrected code)

3.1 Si 2×2×2, 500 bands ($N_p \approx 4483$, $\Pi \approx 177$ GB)

Seconds per SCF iteration	CPU		GPU		
	1n·96r	2n·192r	1n·8×H100	2n·8×A100	4n·16×A100
update_w ($\Pi \rightarrow W$)	343	235	185	162	91
eval Π (RPA)	32	21	21	20	11
$\tau \leftrightarrow \omega$ FTs (redistributes)	37	24	133	111	60
W-Dyson solve (LU + gemms)	273	191	11.6	17	12.8
W D→H mirror + allocs	1	2	17.5	13.4	6.8
THC-GW Σ	88	55	45	50	33
$G_{ij} \rightarrow G_{uv}$ (host-bound gemm)	27	19	17.5	18	9.7
Wall, 3 iters + setup	1423	1018	941	1009	632
host-only phases (Dyson/ μ , mix, write)	109	132	215	337	240

Table 3: kp222/500b phase breakdown. Single node, *no network*: 8×H100 vs. 96-core Genoa is 1.5× wall while its W-solve alone is 23× faster — the 133 s of “communication” is pure PCIe staging + intra-node host MPI. The staging, not the wire, is the problem.

3.2 Where a GPU iteration goes (2n·8×A100, 215 s solver time)

Component	s/iter	class
$\tau \leftrightarrow \omega$ redistributes (host-staged)	111	transfer waste
Necessary device compute (W solve, Π , Σ gemms, FTs)	≈50	device compute
$G_{ij} \rightarrow G_{uv}$ host gemm + per- (t, s, k) H2D	18	host-bound compute
W D→H mirror + 23 GB device allocations	13	transfer/alloc waste
Σ host reduce + shm accumulate	≈12	host-bound
Z re-gather (12.5 s at 16r; 187 s/iter at kp444)	4	transfer waste

Table 4: Waste decomposition. In addition, each run carries ≈200–300 s of host-only SCF machinery (Dyson G -solve, μ search, eigenspectra, mixing, rank-0 serial HDF5 of a 35 GB checkpoint) that a 192-rank CPU job hides but an 8-rank GPU job cannot.

3.3 Anatomy of one redistribute

Each `tau_to_w/w_to_tau` performs two full redistributions of the Π/W tensor (t -distributed $\rightarrow PQ$ -distributed $\rightarrow \omega$ -distributed); four per iteration. For device arrays, `redistribute_alltoallv` (`nda_utils.hpp:715`) does, per peer: strided pack into a fresh device temporary (`raw cudaMalloc`) $\rightarrow$ bulk D2H into a host vector $\rightarrow$ host `MPI_Alltoallv` $\rightarrow$ H2D $\rightarrow$ device read-modify-write unpack. That is four PCIe crossings of ~22 GB/rank per transform pair, $O(n_{\text{ranks}})$ device allocations each, and ~4 s per 23 GB buffer allocation. The May CUDA-aware attempt (0e8fc80) was reverted for the right symptom and the wrong reason: it kept *unchunked* full-size device staging buffers, doubling device memory $\rightarrow$ OOM. The pattern is sound; it needs chunking and a pool.

3.4 Production scale: Si 4×4×4, 500 bands ($\Pi \approx 1.42$ TB)

CPU reference (8 nodes, 192 ranks, 24/node $\times$ 4 threads): **update_w** 1427 s/iter (W-solve 944, **Z re-gather** 187, FTs 134), **THC-GW** 383 s/iter ($G_{ij} \rightarrow G_{uv}$ 118, $\bar{\Sigma}_{uv} \rightarrow \Sigma_{ij}$ 146); wall 4213 s for 2

iterations. Roofline: the W solve is 12.7 Pflop/iter $\rightarrow \approx 82$ s on 64 A100. Every overhead term scales *worse* than compute at kp444 (redistribute volume $\times 8$; Z-gather $\times 8$ with a full- N_p^2 allreduce per q -point). GPU runs at this scale (64 $\times$ A100, 48 $\times$ H200, on pre-built ERIs) are queued as before/after validation; they do not change the plan.

4 Improvement plan

Ordered by impact $\div$ risk. Gains quoted against the 2n-8 $\times$ A100 kp222/500b iteration (215 s solver + ~ 110 s/iter-equivalent host phases); P1/P4/P6 gains are larger at kp444.

P1 — Device-direct chunked redistribute (NCCL). **−95–100 s/iter.**

(Effort M, risk M.) DEVICE specialization of `redistribute_alltoallv`: pack sub-blocks on device into fixed-size per-peer chunks (1–2 GB workspace from the P2 pool), exchange via `ncclGroupStart/ncclSend/ncclRecv` on the existing `mpi3::nccl` device communicator (already used by the ERI builder), unpack on device. Chunking bounds the extra device memory — fixing exactly what killed the May attempt. Remove the unpack read-modify-write. Fallback: chunked CUDA-aware `MPI_Alltoallv` behind the same interface. At NVLink + GPUDirect rates, 4×22 GB/rank moves in 8–15 s vs. today’s 111 s. Validate with the shared-ERI harness (energies to 10^{-13}).

P2 — Activate `nda`’s existing pooled allocator (no new allocator). **−12–15 s/iter; removes OOM fragility.**

(Effort S, risk L.) Every device scratch array is currently a raw `cudaMalloc/cudaFree` (`nda::cuarray = heap<Device> = mallocator`); a 23 GB allocation costs ~ 4 s and each redistribute makes $O(n_{\text{ranks}})$ of them. CoQuí already exposes the remedy — `memory::buffered_array<MEM, ...>`, backed by `static_fallback<dynamic_bucket<Device>` (`configuration.hpp:169-211`) — a first-fit pool with free-list coalescing, a process-wide static instance, and graceful fallback to `mallocator` on a pool miss. *No new allocator is needed.* **Why it isn’t already helping:** the pool is inert — `dynamic_bucket`’s default capacity is 8000 bytes and nothing in the tree ever calls `resize()`, so even the four existing `buffered_array` uses (cuSPARSE glue) fall straight through to `cudaMalloc`. Work: (i) size the pool once at startup via `get_primary()->resize()` from a config knob (fraction of free VRAM); (ii) route the hot paths — either flip `device_array` (`configuration.hpp:105`) to the buffered handle `detail::buffered_handle_t<DEVICE_MEMORY>`, which pools every `memory::array<DEVICE_MEMORY, ...>` with no call-site type churn, or opt in per hot path to keep the giant per-iteration darrays on the plain heap. **Constraints:** `resize()` throws if any block is live (size before the first device allocation, never mid-run); first-fit with coalescing has no defragmentation, so mixing 20 GB darrays with thousands of small scratch blocks can fragment the pool and silently revert to `cudaMalloc` — log capacity and `maximum_memory()` to catch it; the static pool is neither thread-safe nor stream-ordered, which is fine today (one rank per GPU, single stream) but would require per-stream pools or `cudaMallocAsync` if P1’s chunking later overlaps pack/send/unpack across streams (the one case where `cudaMallocAsync` retains an edge). Note also that reserving a large pool shrinks the `cudaMemGetInfo` free memory the ISDF builder reads to choose its blocking (`thc_aux.icc:1189`): size the pool after ERI construction, or make that query pool-aware.

P3 — Device-resident basis transforms. **−25–30 s/iter.**

(Effort M, risk L.) In `_primary_to_aux_impl` the first gemm ($X \cdot G$) runs on the host (one core

per rank) with a small H2D per (t, s, k) — which is why $G_{ij} \rightarrow G_{uv}$ costs the same 18 s on 8 GPUs as on 192 CPU cores. Stage G (4.4 GB at kp222; chunked over k/t at kp444) and the X collocation blocks to device once per iteration and run both gemms there. Same change serves Π and Σ ; drop the per- i host accumulate on the output side.

P4 — Cache iteration-invariant THC data. -5 s/iter (kp222); -180 s/iter (kp444).

(Effort S, risk L.) `thc_reader_t::Z()` re-gathers each q block every iteration with two scalar all-reduces, a barrier-bearing gather, a full- N_p^2 allreduce (only the local block is kept), and a fresh H2D. Z never changes during SCF: cache the device slice per local q (LRU if tight). Also cache the small IAFT transform matrices (rebuilt + re-uploaded per call) and the $k \leftrightarrow R$ coefficients (recomputed up to $8\times$ /iter, each with a global barrier + internode allreduce).

P5 — Lazy W host mirror. -13 s/iter; -180 GB host RAM/node.

(Effort S, risk L.) `update_w` keeps `dW_qtPQ_dev` for the GPU Σ path yet still materializes the full host mirror each iteration for consumers absent in a pure scGW run. Make it lazy; batch the $(t, q) \rightarrow (q, t)$ transpose into one kernel.

P6 — Device Dyson stage; eigenspectra memory; μ cap. $-60-80$ s/iter-equiv.

(Effort M/L, risk M.) `simple_dyson` has no MEM template: $n_\omega n_k$ inversions of n_b^2 matrices, the μ search, and eigenspectra run on host across 8–64 ranks. Add a DEVICE path for the Dyson loop (batched `getrf/getrs` over ω), distribute/chunk the per-rank $\Sigma(\omega)$ buffer (B5), cap and log the μ walk (B4). The host `geev` is cheap once the redundant allocation is gone.

P7 — LU-solve instead of explicit inverse in the W solve. $-6-7$ s/iter.

(Effort S, risk L.) $W = [1 - Z\Pi]^{-1}Z - Z$ currently does `getrf+getri+gemm`; solving $[1 - Z\Pi]W' = Z$ with `getrf+getrs` removes $\sim 40\%$ of the flops and one gemm. When the (ω, q) block is rank-local (the common case) call cuSOLVER directly, skipping the per-call SLATE object construction (7 per (ω, q) pair). Remove the diagnostic `device_sync()` serialization in production builds.

P8 — Fix the device ISDF builder ≥ 16 ranks (B3); Σ divergence correction.

(Effort M, risk M.) Trace the `Zqur`-stage device pointer that reaches a host-only MPI path; route through NCCL like the builder’s other collectives, or stage that single exchange. Until then the `save=/load` workaround stands. `Sigma_div_correction` is host-only scalar $O(n_s n_k n_b^2 N_p)$ per iteration — port or thread before kp444-scale GPU production.

P9 — Write path. $-20-30$ s/iter-equiv.

(Effort S, risk L.) Rank-0 serial HDF5 of ~ 35 GB per iteration (85–90 s per kp222 GPU run). Parallel HDF5 or an async writer thread; optionally checkpoint every N iterations.

4.1 Projected outcome (kp222/500b, 2 nodes · $8\times$ A100)

Per SCF iteration	today	after P1–P7	change
<code>update_w</code>	162 s	≈ 38 s	redistributes $111 \rightarrow 12$, W-solve $17 \rightarrow 11$, mirror/allocs $13 \rightarrow 1$, Π $20 \rightarrow 14$
THC-GW Σ	50 s	≈ 15 s	$G_{ij} \rightarrow G_{uv}$ $18 \rightarrow 2$; host reduce partially remains
Host phases (/iter-equiv.)	~ 110 s	≈ 45 s	device Dyson, capped μ , async write
Total per iteration	≈ 320 s	≈ 100 s	vs. CPU 2n ≈ 340 s $\Rightarrow \approx 3.5\times$ wall, $\approx 5\times$ solver

At kp444 the removed terms scale with n_q and rank count, so the node-for-node advantage grows; the W-solve roofline alone bounds it at $\approx 11\times$ (16 A100 nodes vs. 16 Genoa nodes). A realistic production target is **5–8 $\times$ node-for-node**, at which point the remaining gap is irreducible communication.

4.2 Suggested sequencing

1. **Week 1:** P2 (activate the nda pool: size at startup, route the device handle) $\rightarrow$ P1 (NCCL redistribute) behind a runtime switch; validate with the correctness harness; re-run the kp222 matrix.
2. **Week 2:** P3 + P4 + P5 (transfer hygiene); B4 μ -walk cap. Re-run kp222; launch kp444 GPU.
3. **Week 3+:** P6 (device Dyson + eigenspectra memory), P7, P9, then P8 (builder fix) to retire the ERI-prebuild workaround.

5 Run inventory

Under `~/ceph/CoQui/GPU_PORT_run/` on rusty:

- `si_kp222_n500_e125/bench_{cpu,gpu}_*` — kp222/500b matrix: CPU 1n/2n; A100 2n/4n; H100 1n/2n (post-fix logs: jobs 6674560–62, 6676687, 6678589–90); `eri_n500.h5` (2.9 GB).
- `si_kp222_n250_e125/`, `si_kp444_n500_e125/` — 250-band single-node pair; kp444 CPU 8n (6671589); `eri_n500.h5` (23 GB); pending GPU: 64 $\times$ A100 (6678594), 48 $\times$ H200 (6678593).
- `correctness/` — shared-ERI validation matrix, CPU 1/2/4/8 ranks + GPU H100 1/4 ranks; all $E = 0.5146018573397(1-4)$ a.u.
- `GPU_PORT/src` — branch `gpu`; binaries `build/cpu`, `build/gpu` (sm80), `build/gpu90` (sm90); fixes 80f4a5d, 765754c pushed to `origin/gpu`.